

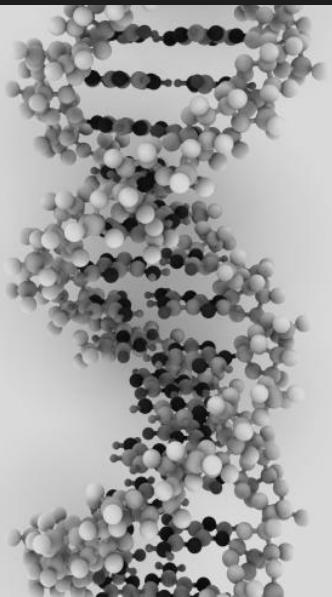


High Performance Spatial Query Processing for Large scale Scientific Data

Ablimit Aji

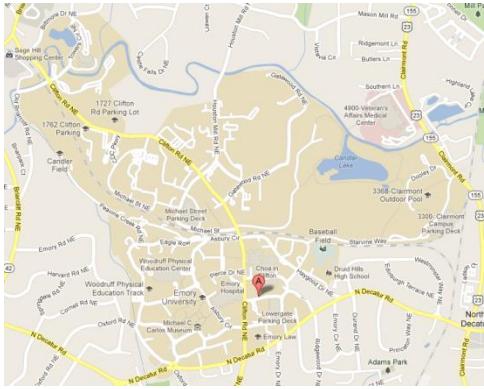
Adviser: Fusheng Wang

Math & CS Seminar
Nov 16, 2012



Spatial Applications

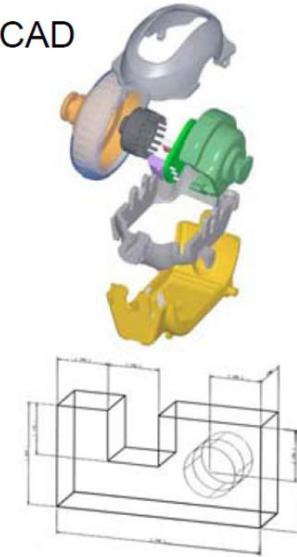
Map



Epidemiology
and health



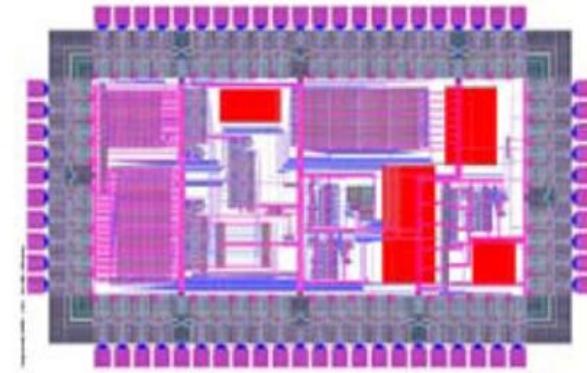
CAD



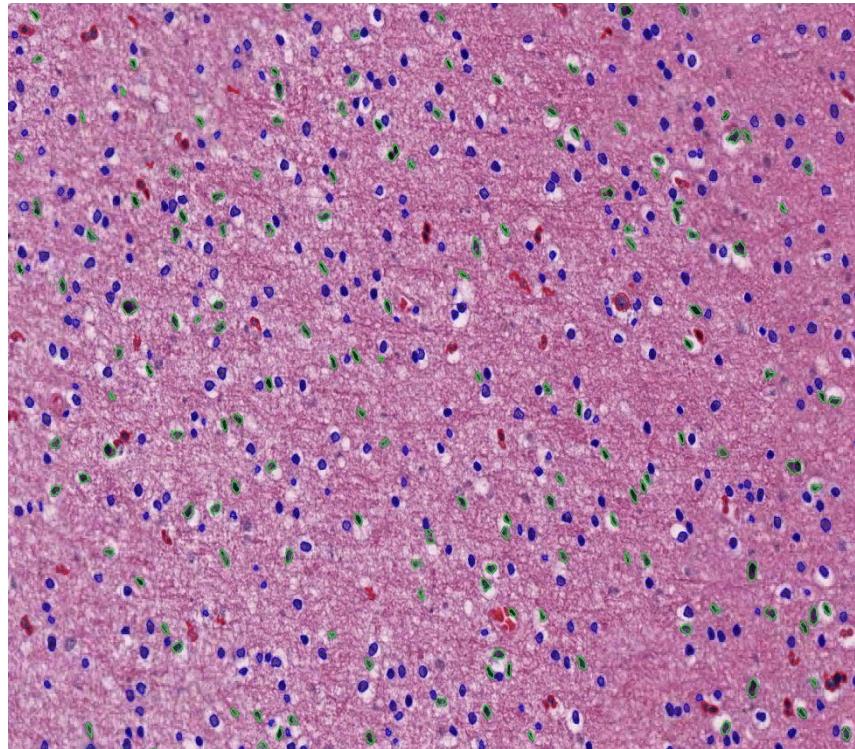
Cosmology



VLSI Design



Big Spatial Data



South America, 2004, orbit 022868, Block B61.
The plume region is manually labeled with green arrows. Yellow arrow is wind direction. Bright white features are clouds. Red spots are MODIS retrieved fire spots.

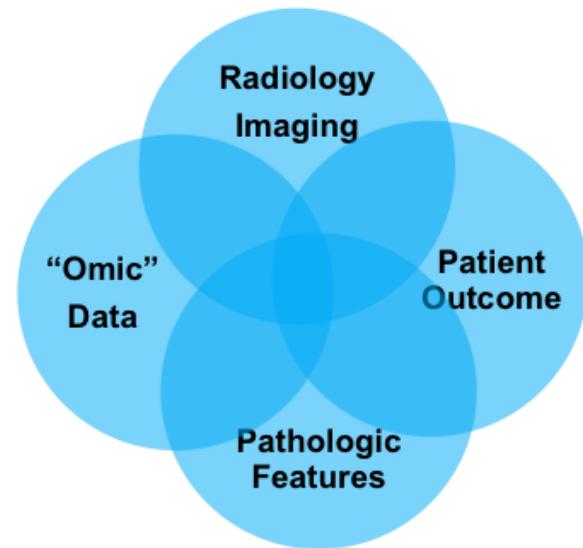
Figure 1. Smoke Plume of Satellite Imagery

Pathology Analytical Imaging

Satellite Imagery & Remote Sensing

Integrative Multi-Scale Biomedical Informatics

- Reproducible anatomic/functional characterization at gross level (Radiology) and fine level (Pathology).
- Integration with multiple types of “omic”s (genomics, proteomics, metabolomics) information.
- Create categories of jointly classified data to describe pathophysiology, predict prognosis and response to treatment.

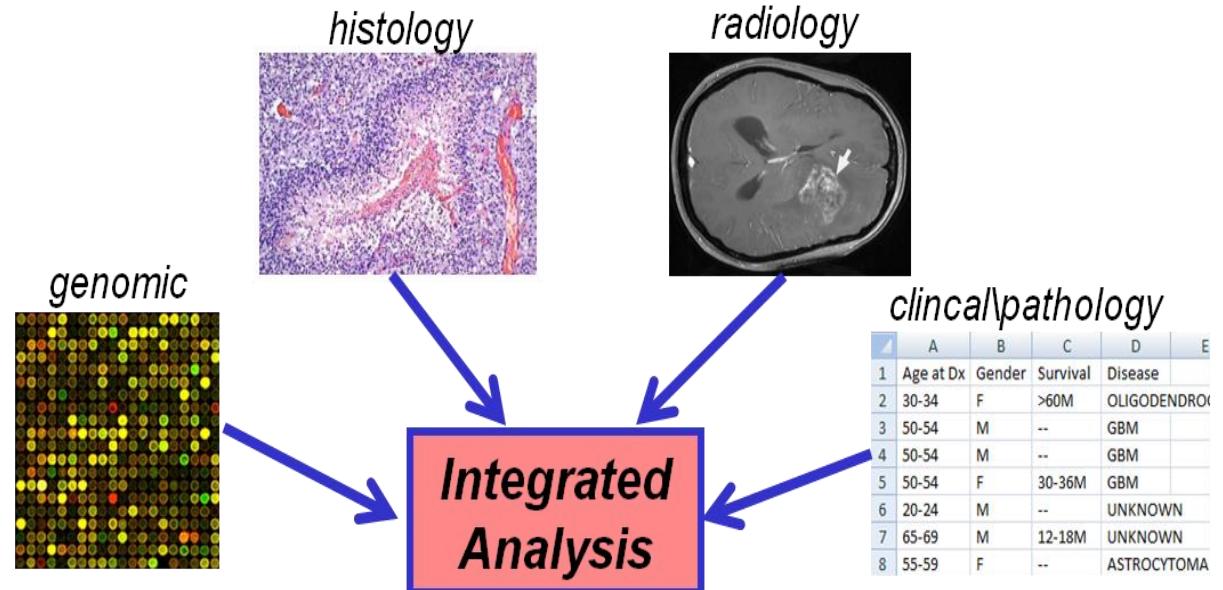


Emory In Silico Brain Tumor Research Center

Integrative in-silico study of GBM brain tumors through digital pathology, radiology, “omics” data, and patient outcome.



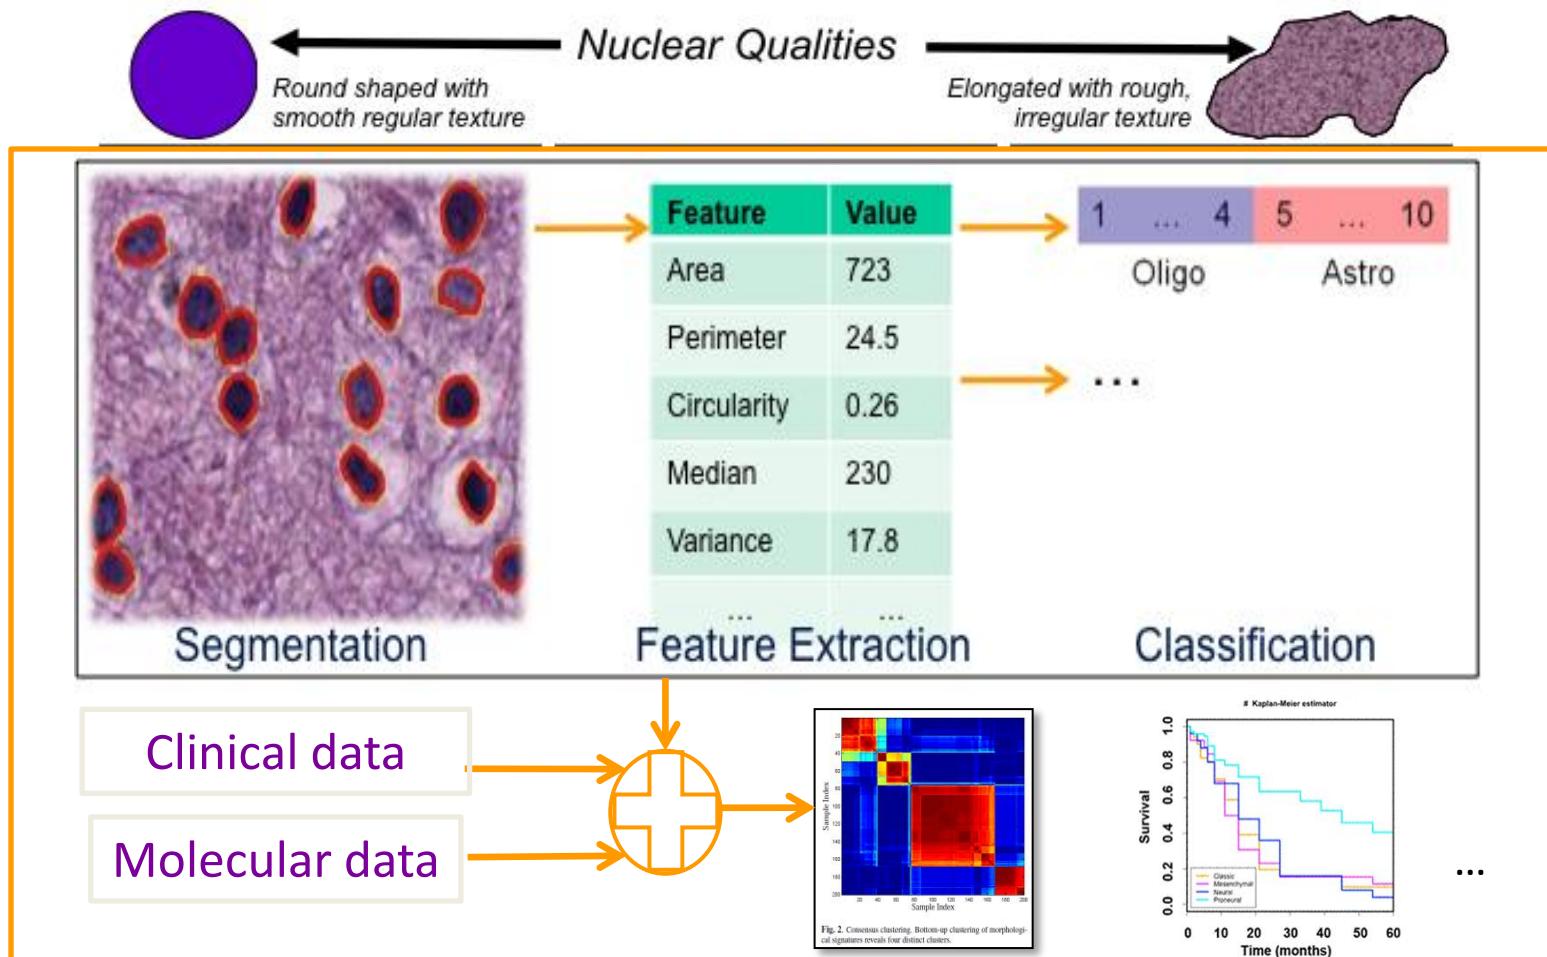
Emory In Silico
Brain Tumor
Research Center



Pathology Imaging

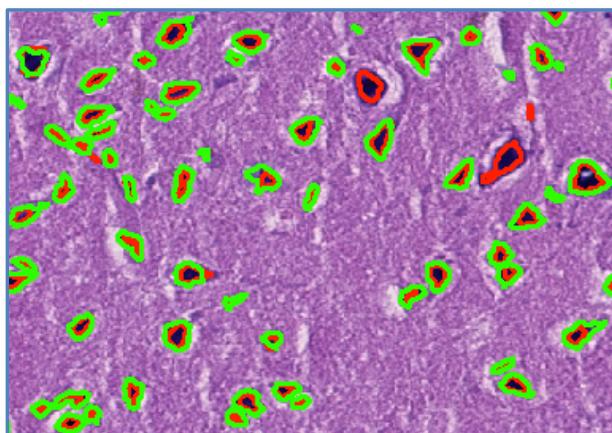
- High-resolution whole slide images provide rich information about morphological and functional characteristics of biological systems. ([Ex1](#))
- Such information has tremendous potential for providing insights regarding the underlying mechanisms of disease onset and progression.

Distinguishing Characteristics in Pathology Images

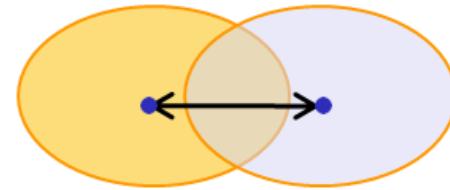


Systematic Image Algorithm Evaluation

- High quality image analysis algorithms are essential to support biomedical research and diagnosis
 - Validate algorithms with human annotations
 - Compare and consolidate different algorithm results
 - Sensitivity study on algorithms' parameters
- Example: What are the distances and overlap ratios between markup boundaries from two algorithms ?



Cross match / join two spatial data sets



Algorithm one v.s. Algorithm two

Spatial Centric Queries

POINT



CONTAINMENT



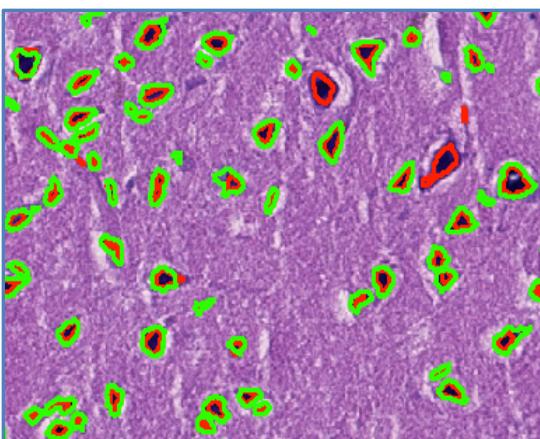
WINDOW



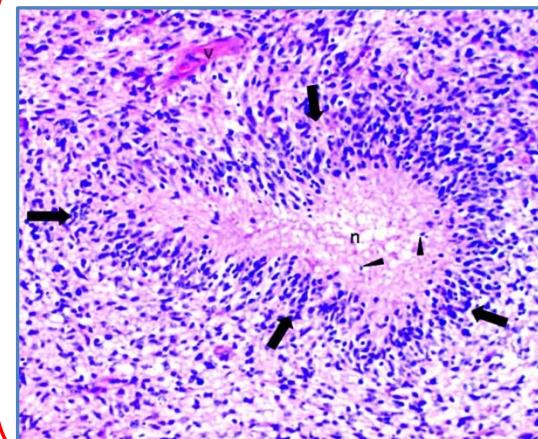
Need: Manage, Query and Compare
Spatially Derived Information



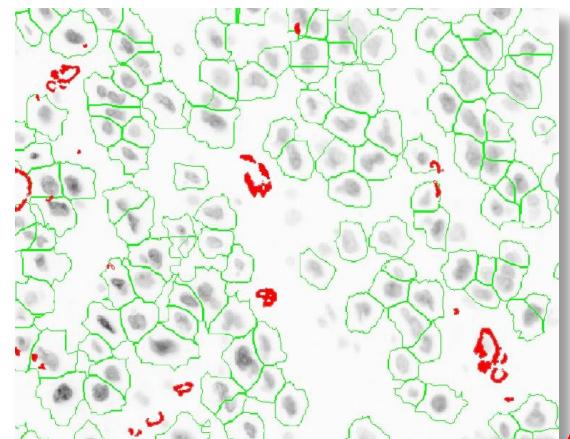
SPATIAL JOIN



DENSITY

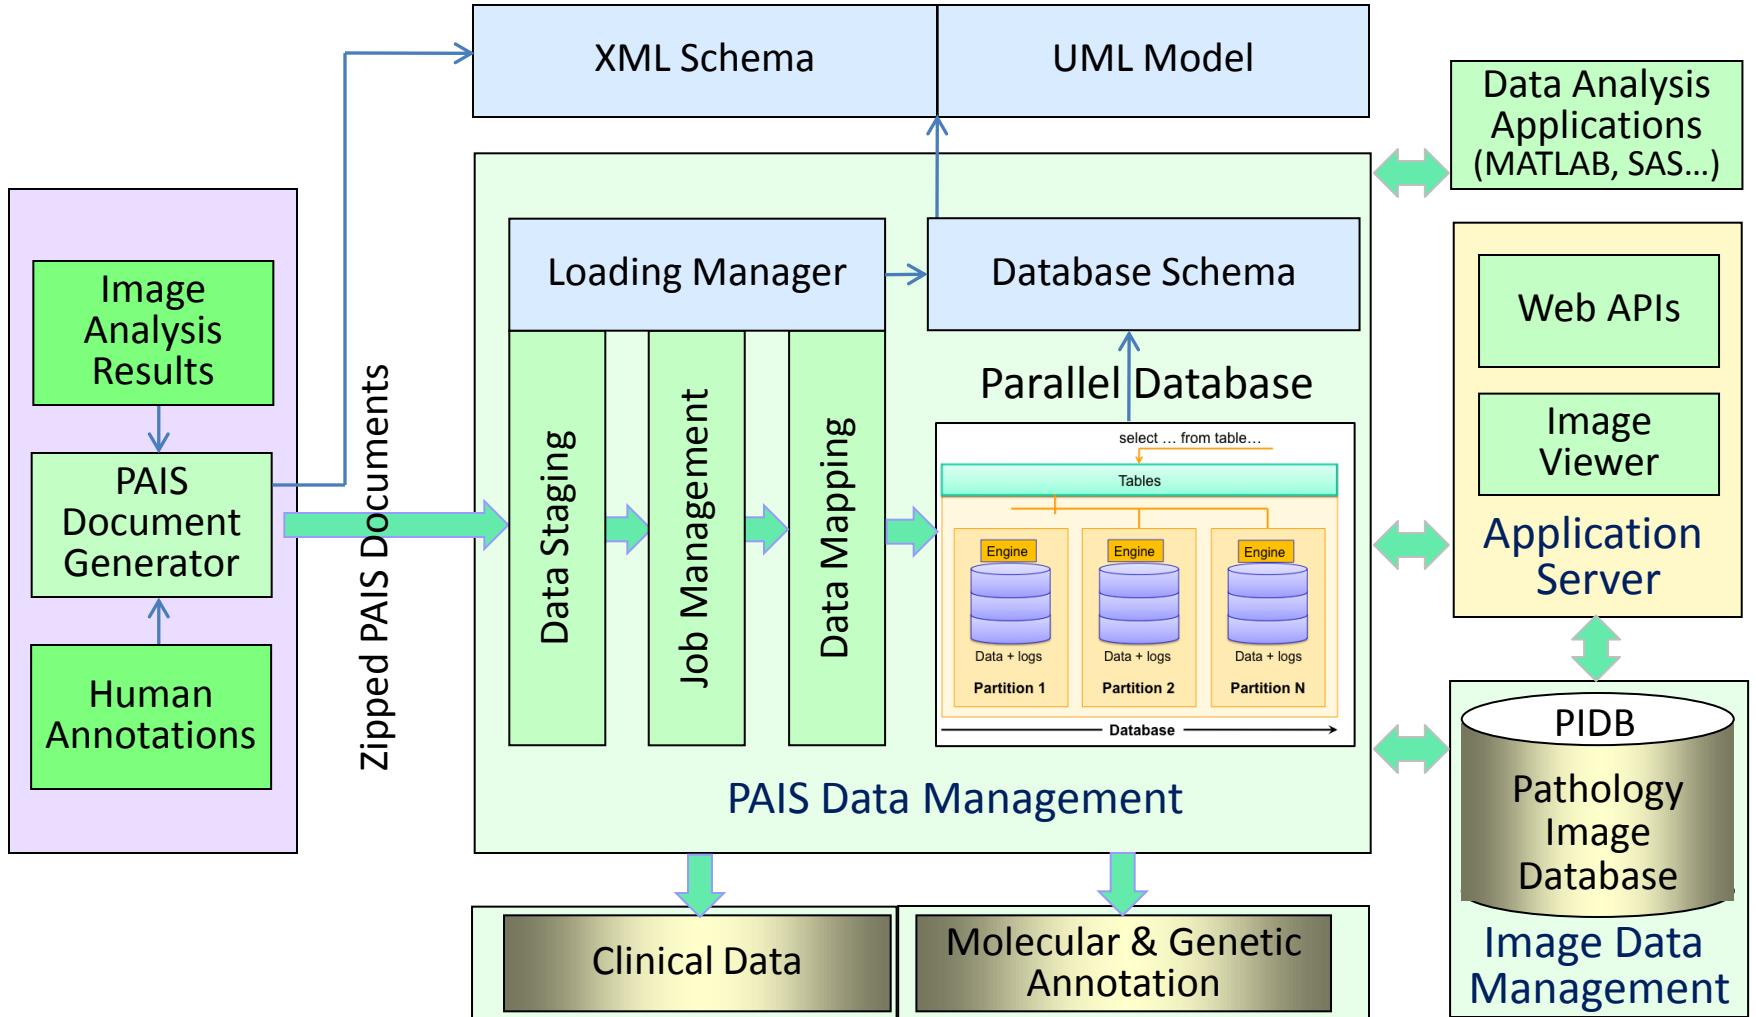


NEAREST NEIGHBOR



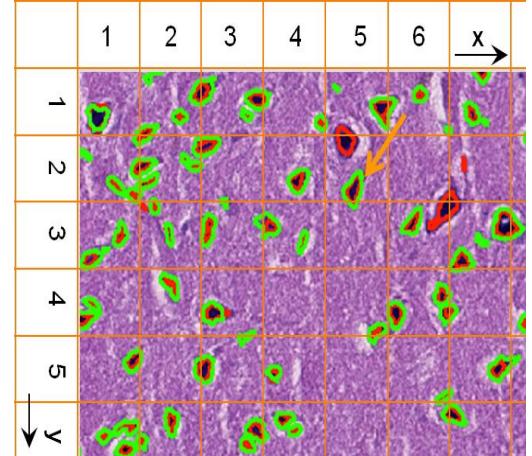
EMORY
UNIVERSITY

PAIS Data Management Architecture



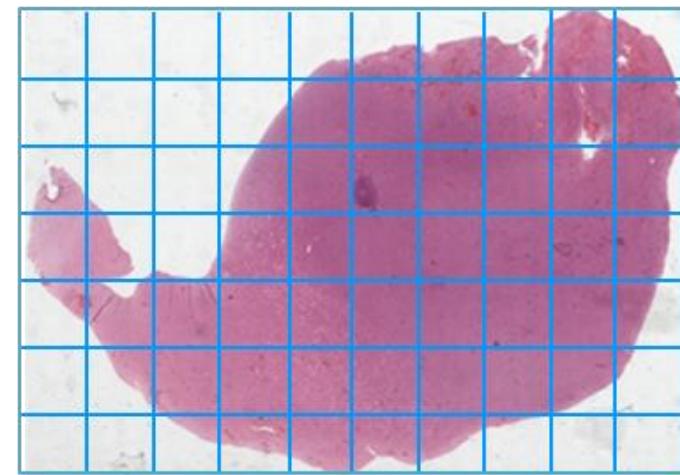
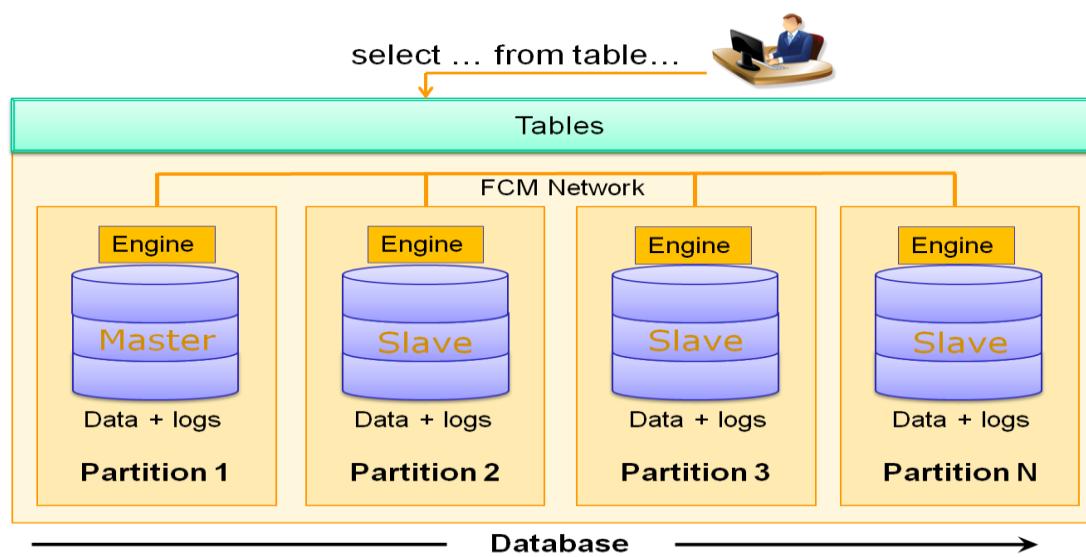
Spatial Database Engine

- Extended on RDBMS to support multi-dimensional spatial data types and queries
 - Extended spatial data types
 - Spatial functions and procedures
 - Spatial access methods
 - space oriented – e.g. Grid-File
 - data oriented – e.g. R*-Tree



Partitioned, Shared-Nothing Parallel DBMS

- Most often I/O is the main bottleneck
- Partition data for parallel data access
- Shared-nothing architecture is widely used in practice as a scalable database solution



Parallel Spatial DB Summary

- Comprehensive data model to represent data and provenance.
- Expressive query interface (Spatial SQL).
- Can be scaled with a partitioned parallel DBMS.
- Can be quickly setup, and support a set of queries for small/medium scale data.

Questions ?

Limitations of Parallel Database Approach

- Data loading takes very long time
- Limited scalability: *possible but with high cost*
- Expensive software/hardware license
- Limited query support
- Maintaining & tuning is complex
 - most often needs a DBA

F. Wang *et al.*, High Performance Analytical Pathology Imaging Database for Algorithm Evaluation, In MICCAI-DCI 2011

Large Volumes of Derived Data

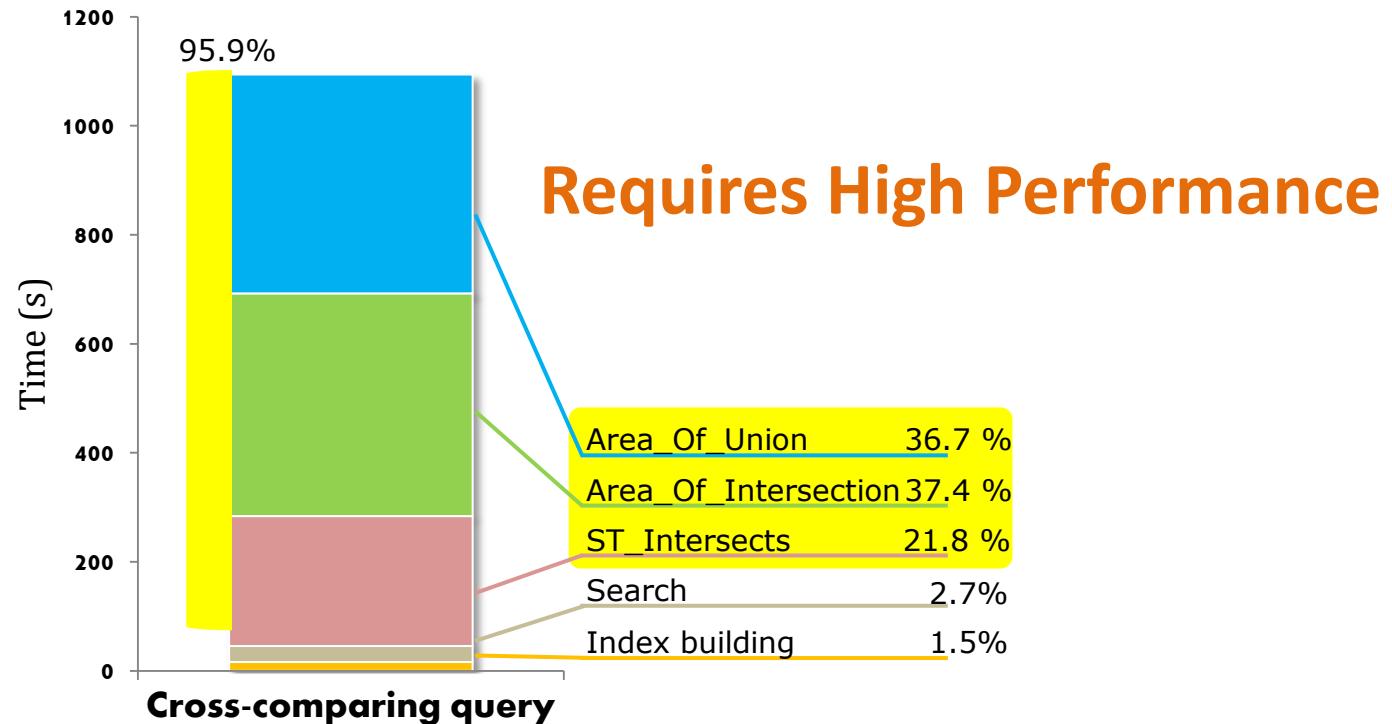
- 10 billion pixels for each WSI: $10^5 \times 10^5$
- 1M derived spatial objects, 100 M features per image
- Thousands of patients, tens of slides /patient
- Many algorithms with varying parameters
- Moderate size healthcare operations can generate Terabytes of data per day easily
- If full WSI would be adopted for clinical environment, petabytes of data per day is likely

Demands High-Throughput

High Complexity of Geometric Computation

```
SELECT AVG(ratio) FROM (
    SELECT
        ST_Area(ST_Intersection(p.the_geom, q.the_geom)) /
        ST_Area(ST_Union(p.the_geom, q.the_geom)) AS ratio
    FROM    oligoastroiii_1_1 AS p, oligoastroiii_1_2 AS q
    WHERE   ST_Intersects(p.the_geom, q.the_geom));

```

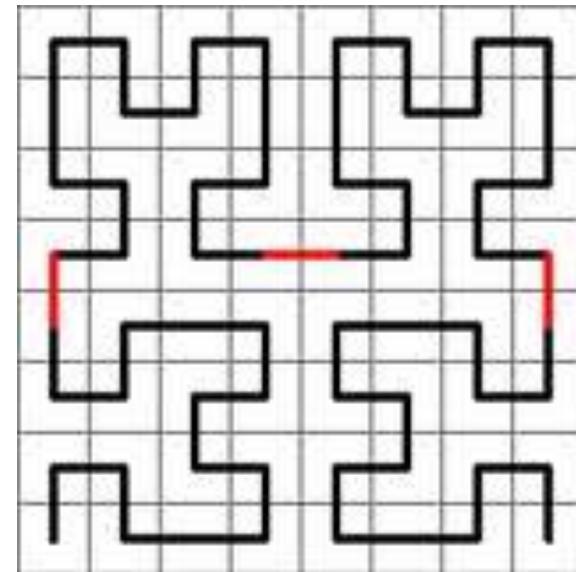
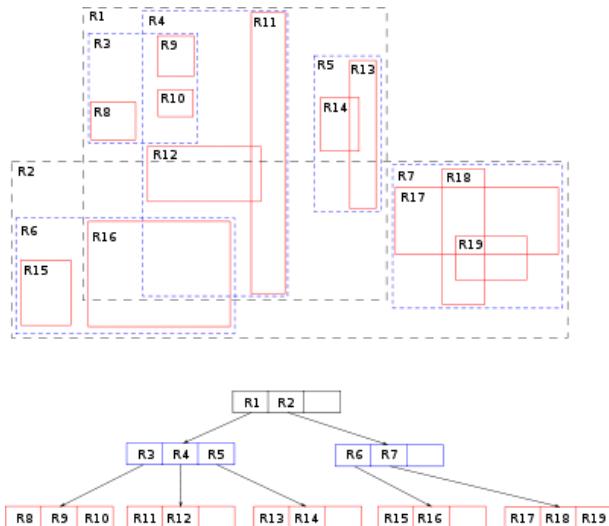


High Performance Queries with MapReduce

- MapReduce is a parallel computing framework widely used for large-scale data analysis and queries
 - Widely used in major internet applications
open source version: [Hadoop](#)
 - Two simple UDFs for data processing : *map* & *reduce*
 - Very easy to setup develop scalable applications
 - Implicit parallelization by data partitioning
- Our approach:
 - Take advantage of MapReduce to run queries
 - Native spatial query support with a spatial engine

Major Gaps

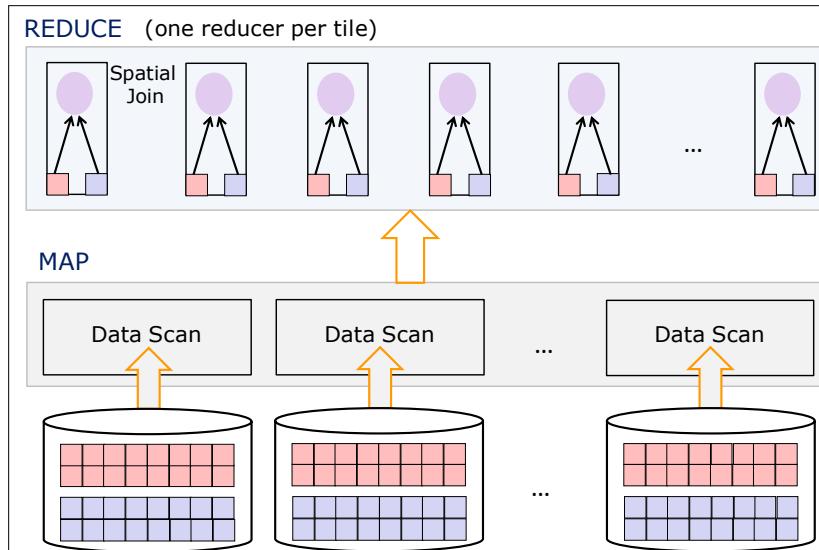
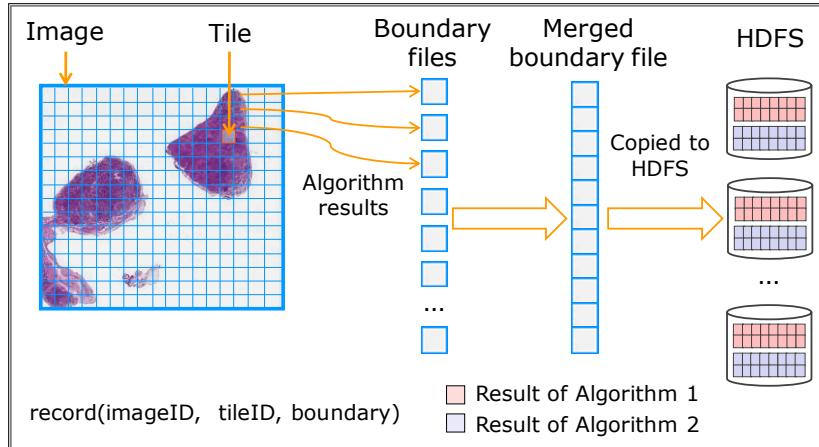
- Spatial Data Partitioning and Index Support
 - Spatial data is multidimensional
 - Hard to preserve data locality
 - Supporting and managing multidimensional index on MapReduce is not easy



Major Gaps

- MapReduce only provides map() and reduce()
 - Need to map spatial query processing algorithms to **map** and **reduce** environment
 - Need to revise and rethink spatial query processing
- An easy to use query interface is required
 - MapReduce == assembly language for MPP
 - Industrial examples – Hive, Pig, Scope ..
 - e.g. in Facebook, 70% of MapReduce workload are submitted through Hive

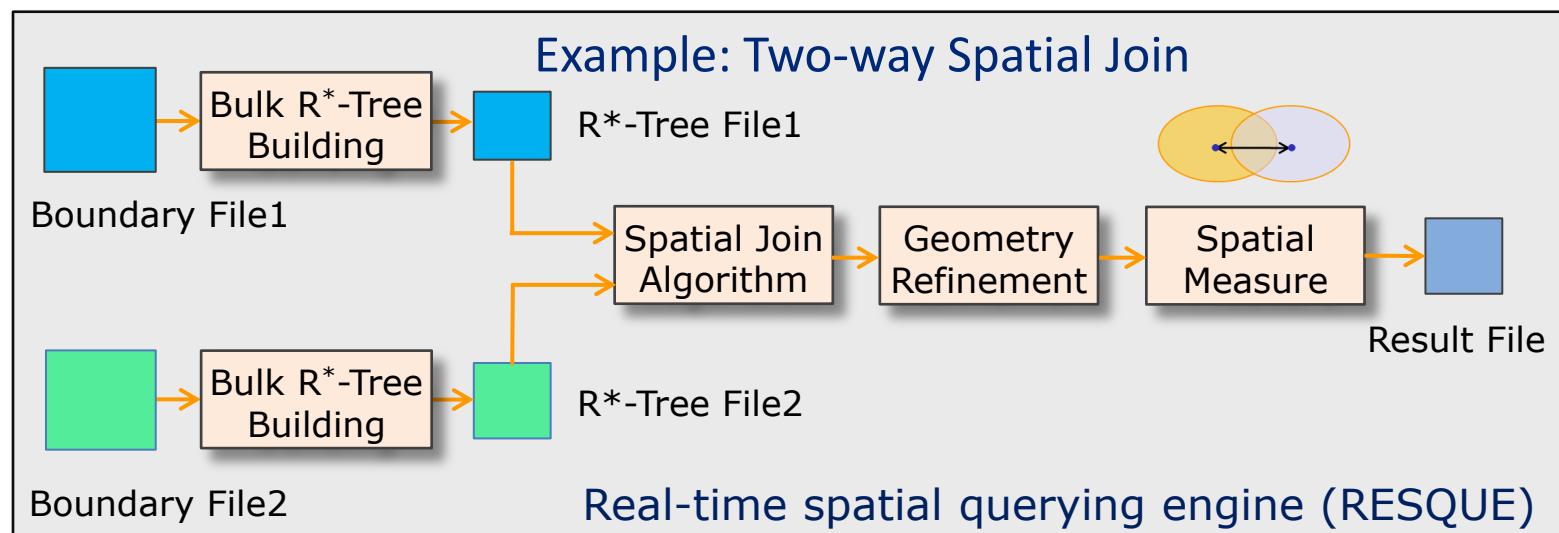
Spatial Join Processing with MapReduce



- **Staging:**
 - images are tiled into regular grids
 - tiles redistributed across HDFS as file blocks
 - small tiles are merged and metadata added into records
- **Map:**
 - identify records of same tiles to form tasks
- **Reduce:**
 - execute queries with the spatial query engine
 - aggregate query results

Real-Time Spatial Query Engine (RESQUE)

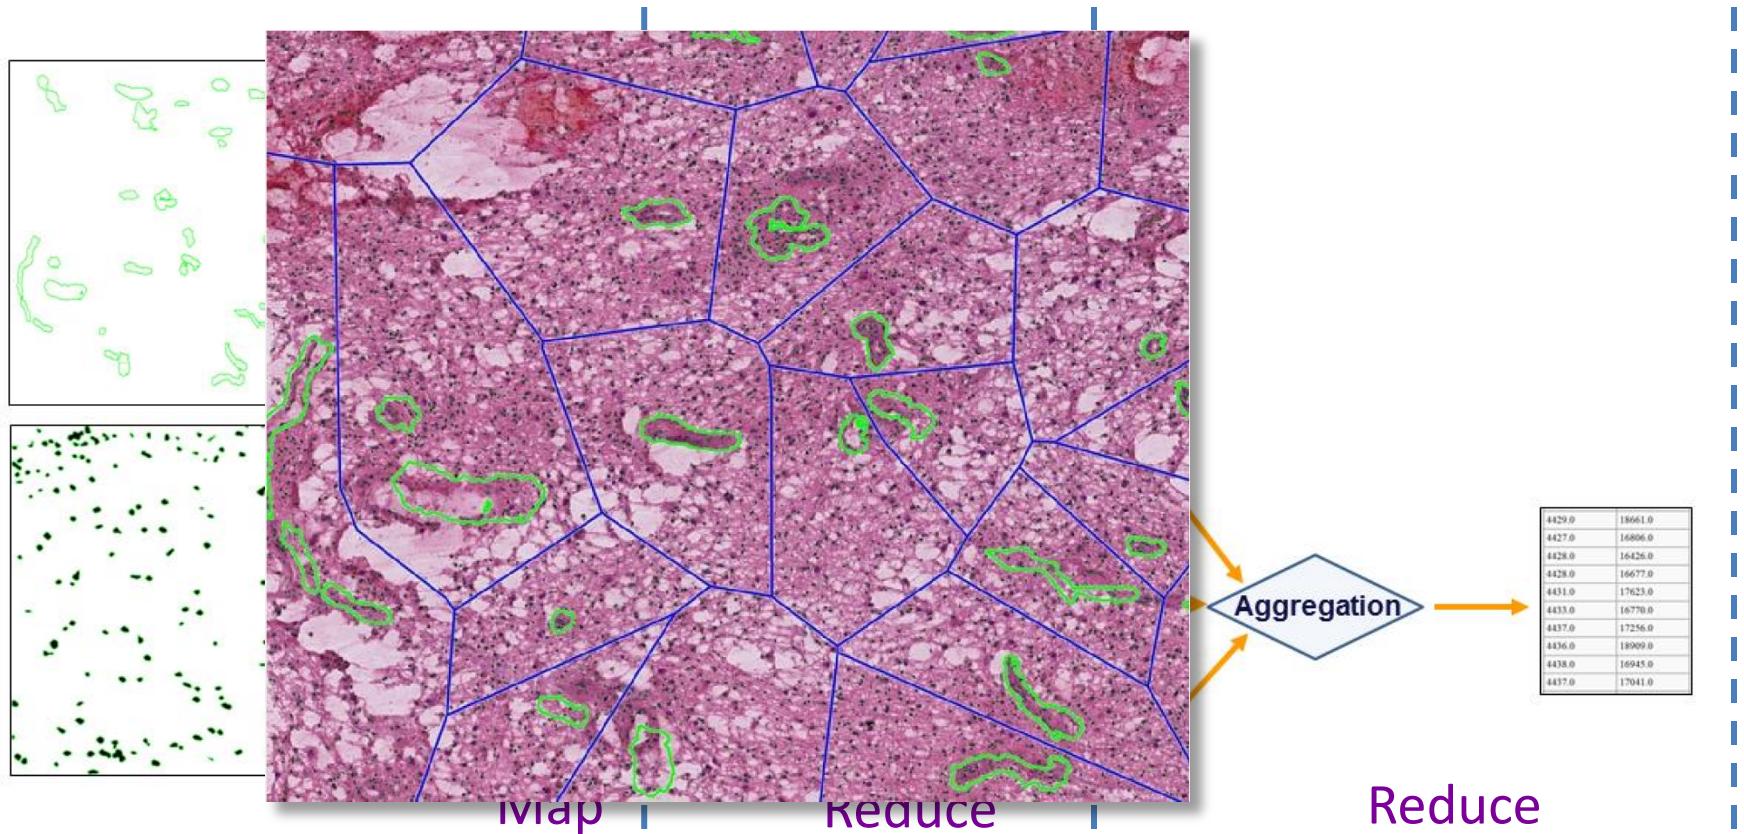
- Index building **on demand** (low overhead)
- Query pipelines to combine multiple steps of query processing
- Decoupled parallel spatial query processing
- Support of spatial access methods and algorithms
 - spatial selection, multi-way spatial join, nearest neighbor, and highest density queries, and extensible for new ones



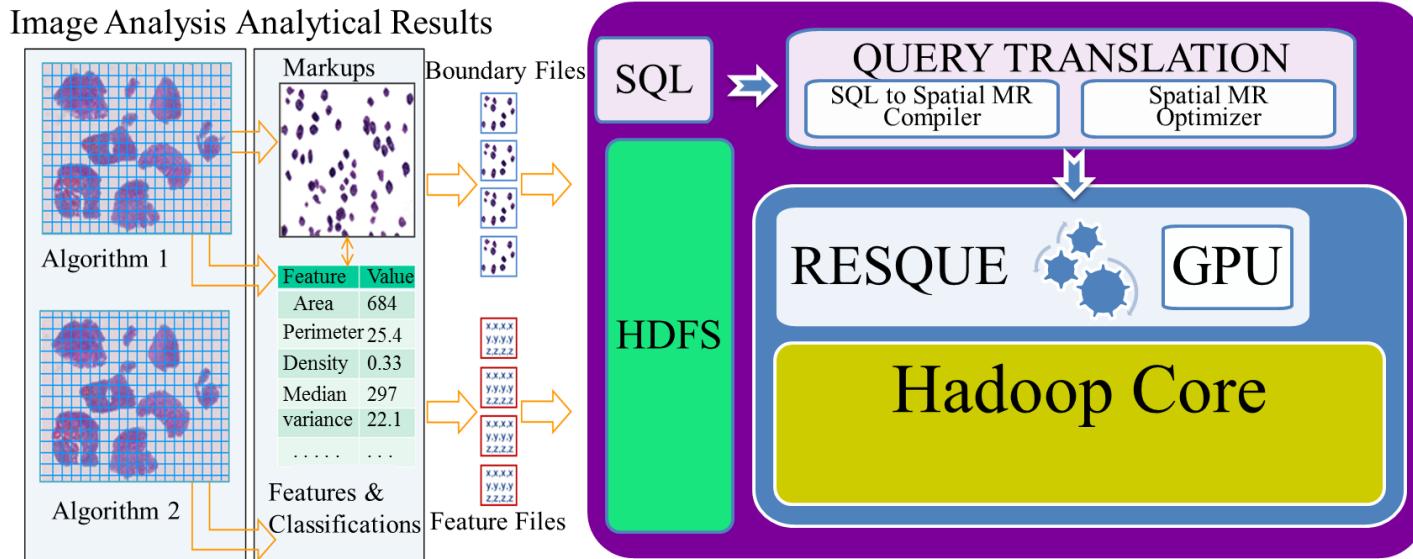
Nearest Neighbor Query Processing Workflow

Query e.g. : for each cell find the closest blood vessel and return distance to that blood vessel.

Access methods can vary (R*-Tree, Voronoi, etc..)



System Architecture – Hadoop-GIS



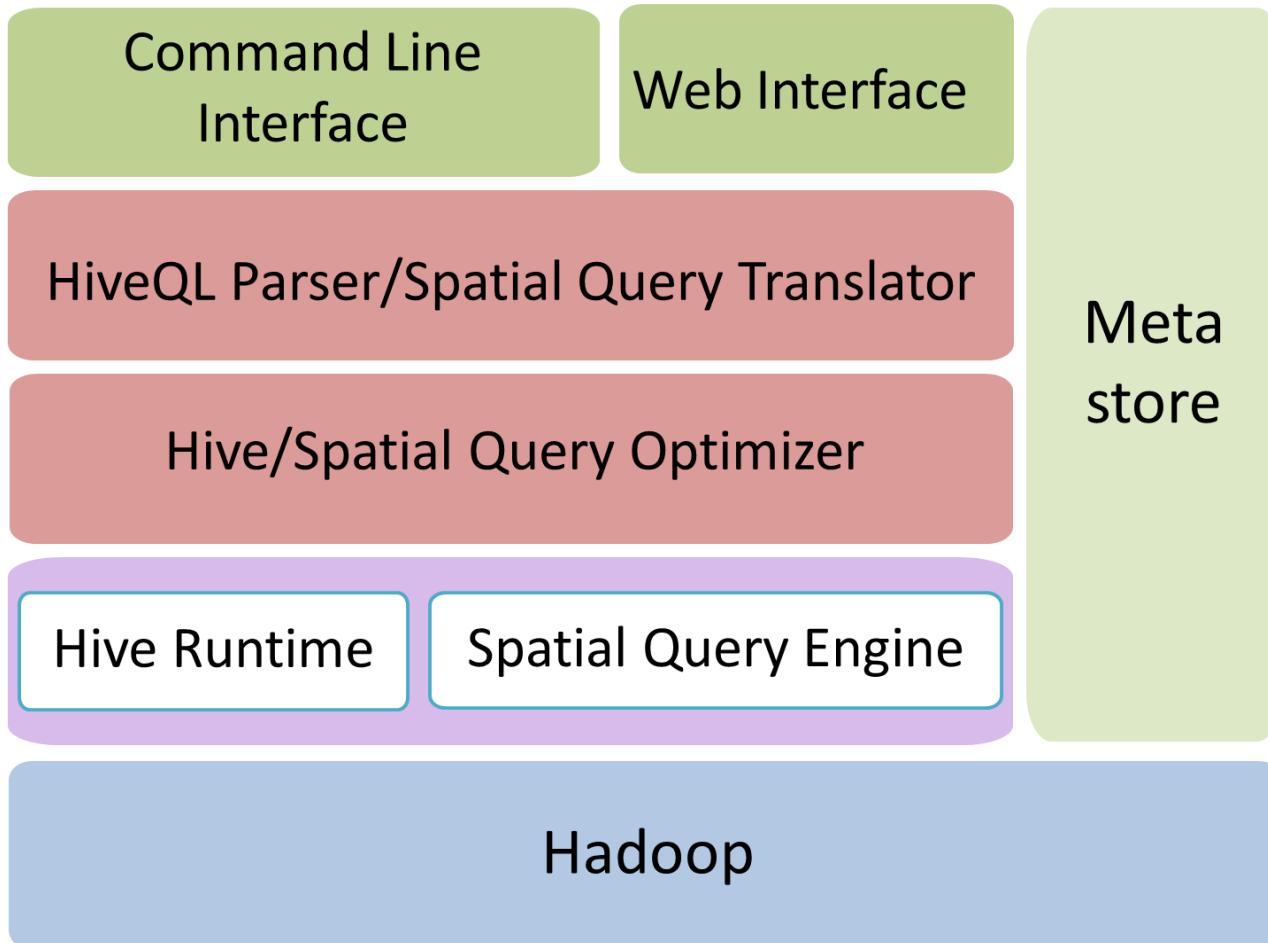
Goal: provide end-to-end support from storage and indexing to query execution (translate MapReduce) for the execution of spatial scientific datasets at big data scales

- Relies on the RESQUE for spatial query processing

Hadoop-GIS

- HiveQL → Hive^{SP}
 - Pros: Widely used as EDW, Tightly integrated with Hadoop
 - Mature query processing pipeline is already in place
 - Cons: Married to Hive, less freedom for optimization
- YSmart → YSmart-Spatial
 - Pros: Good query optimization framework
 - Better control of customized optimization (e.g. GPU)
 - Cons: Not widely used, lots of development

Hadoop-GIS: Hive^{SP}



Declarative Query Interface

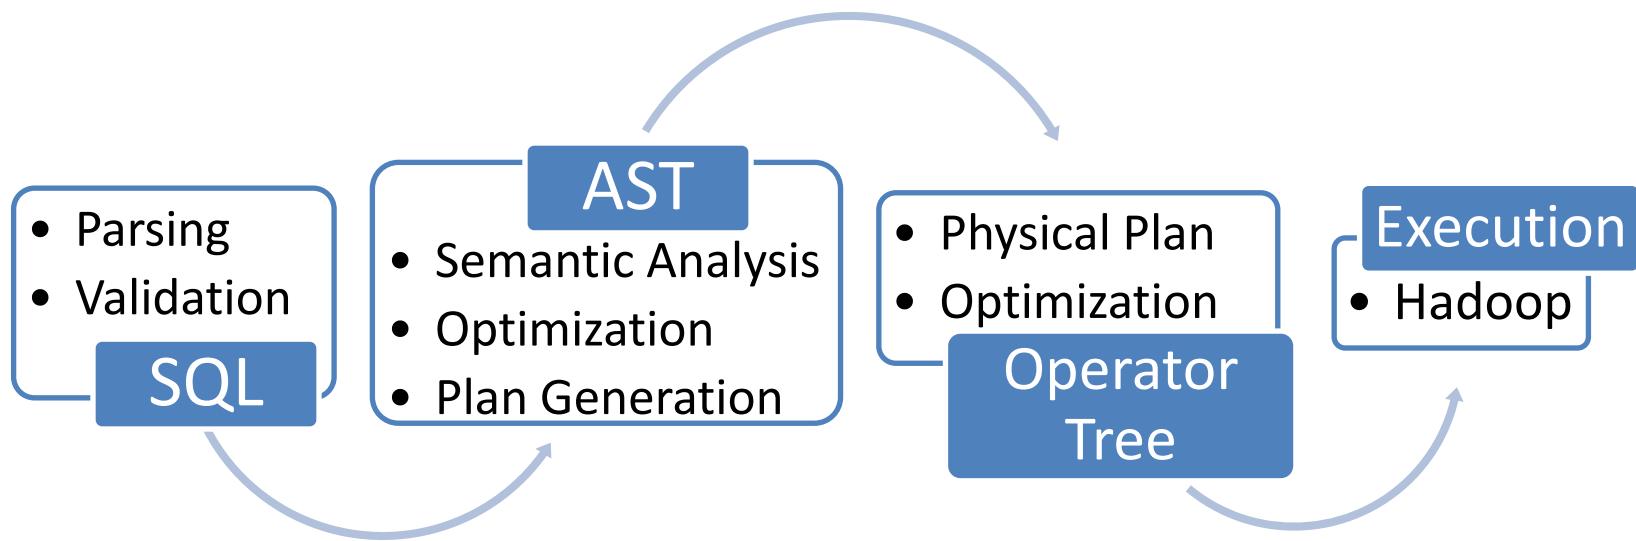
Schema Creation:

```
CREATE TABLE tcga_markups
( markup_id BIGINT, provenance STRING, hand_marked
  BOOLEAN, center ST_POINT, polygon ST_POLYGON)
PARTITIONED BY TILE(polygon ST_POLYGON, 4096, 4096)
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS
TEXTFILE ;
```

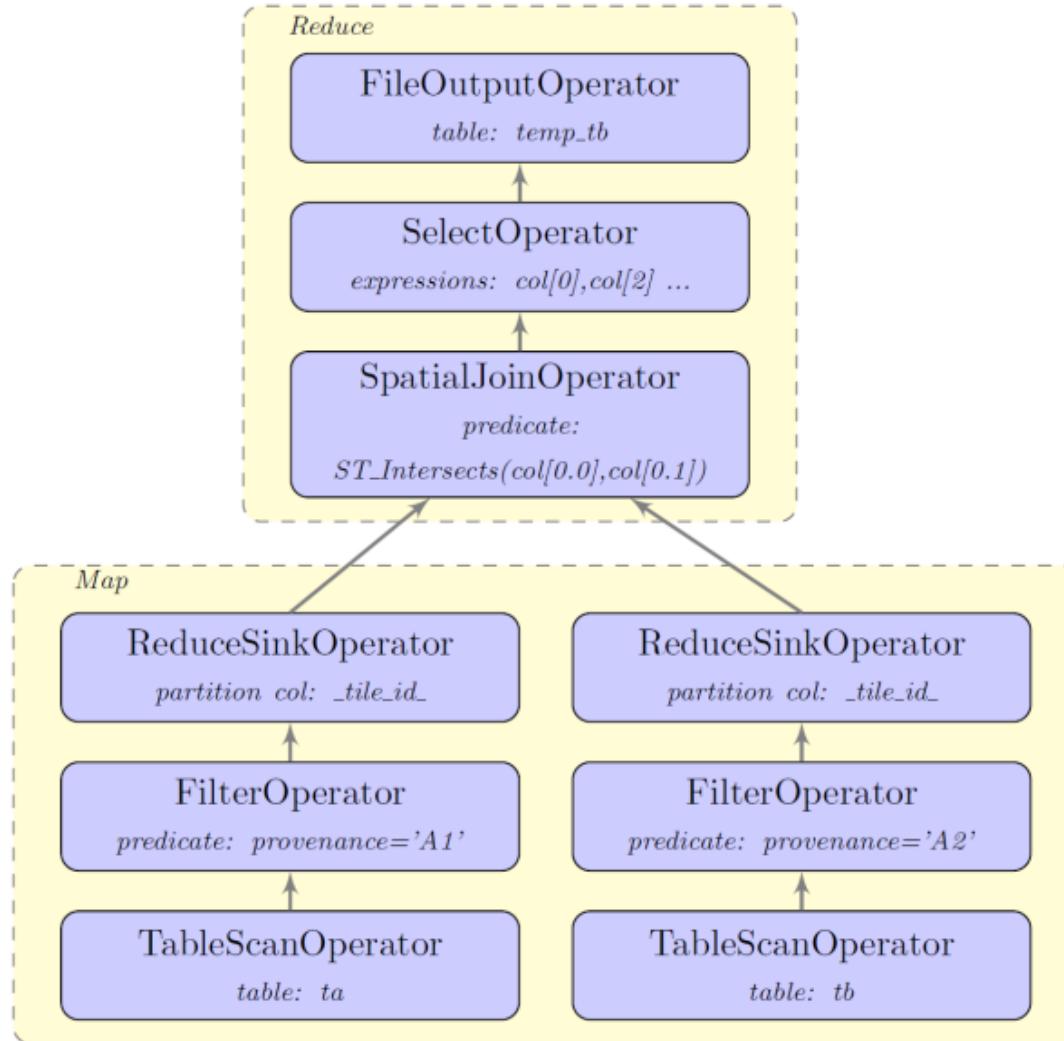
Spatial Join Query:

```
SELECT
  ST_AREA(ST_INTERSECTION(ta.polygon,tb.polygon)) /
  ST_AREA(ST_UNION(ta.polygon,tb.polygon)) AS ratio,
  ST_DISTANCE(ST_CENTROID(tb.polygon),ST_CENTROID(ta.polygon))
AS distance
FROM tcga_markups ta JOIN tcga_markups tb
  (ST_INTERSECTS(ta.polygon, tb.polygon) = TRUE)
WHERE ta.provenance='Algorithm1' AND
tb.provenance='Algorithm2' ;
```

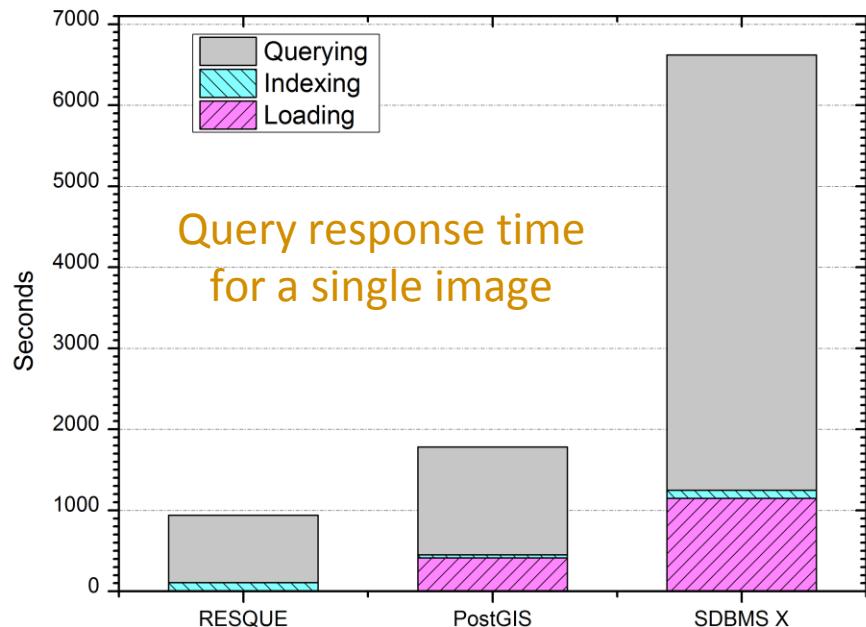
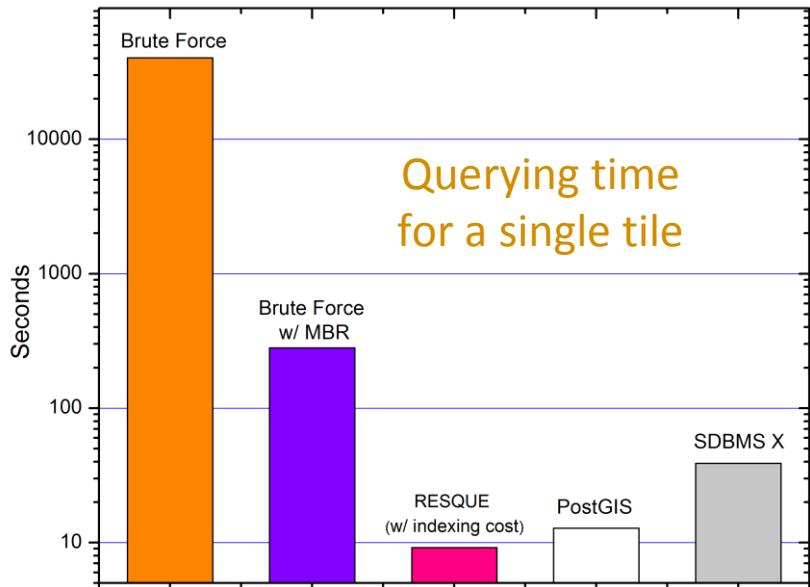
Query Processing Pipeline



Two-way Join Query Plan



Spatial Query Engine Performance



- R*-Tree building $\approx 16\%$ R*-Tree join $\approx 84\%$
- Indexing Scalability: $\text{IndexTime}(\text{Single Image}) \approx \text{IndexTime(Tile)} * \text{tile\#}$
- Storage: apply chain coding at R*-Tree leaves $\approx 42\%$ storage reduction

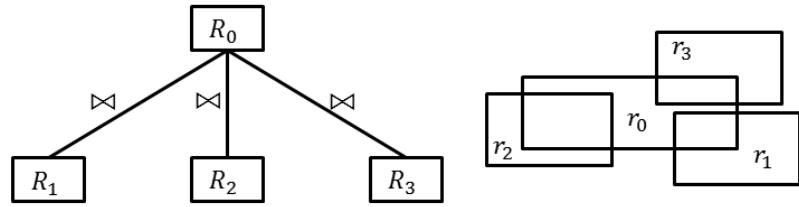
System Scalability Experiments

- Experimental Setup
 - in-house cluster with 10 physical nodes (192 cores)
 - join query: a set of 18 images (0.5 M nuclei/image)
 - 39GBs of data with 6 different results
 - nearest neighbor query: 50 images from TCGA (size)
 - 42 GBs
- Hadoop Setup
 - Cloudera Hadoop 0.20.2
 - Replication factor = 3
 - File split size = 128 MB
 - Individual task tracker memory = 512MB

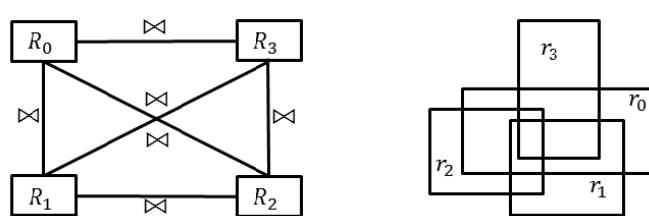
Query Setup

- Multi-way Spatial Join Query

star join



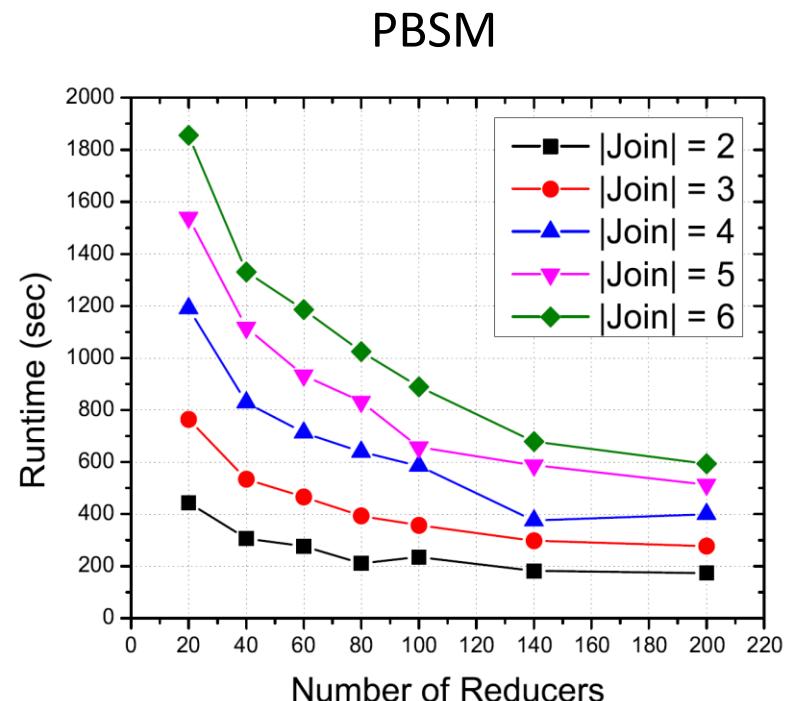
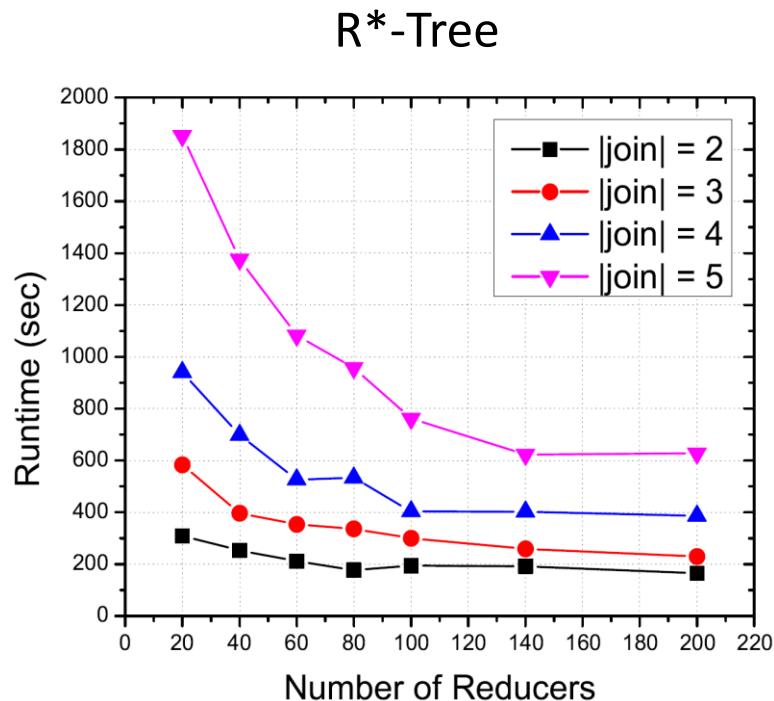
clique join



- Join Algorithms

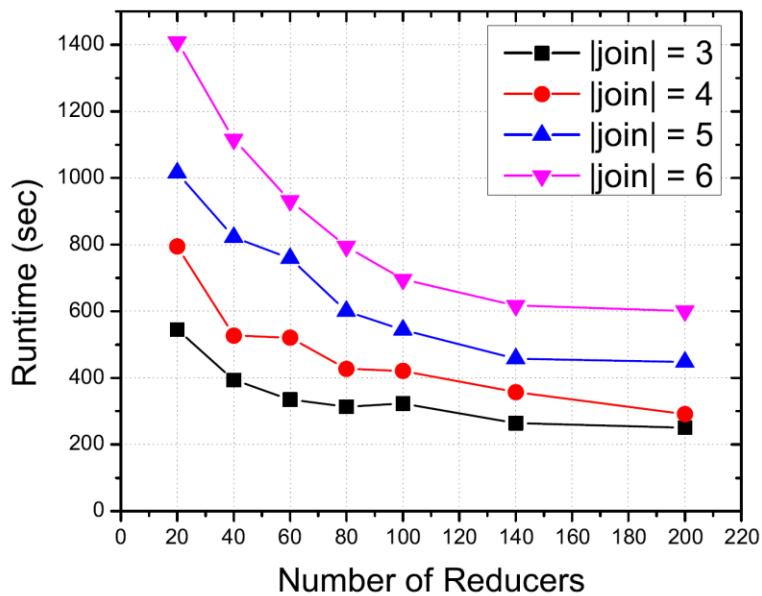
- R*-Tree join (Breadth First Tree Traversal)
- Partition based spatial merge join (PBSM)
 1. refine with approximate processing
 2. filter candidates with precise algorithm

Multi-Way Spatial Join: Star Join

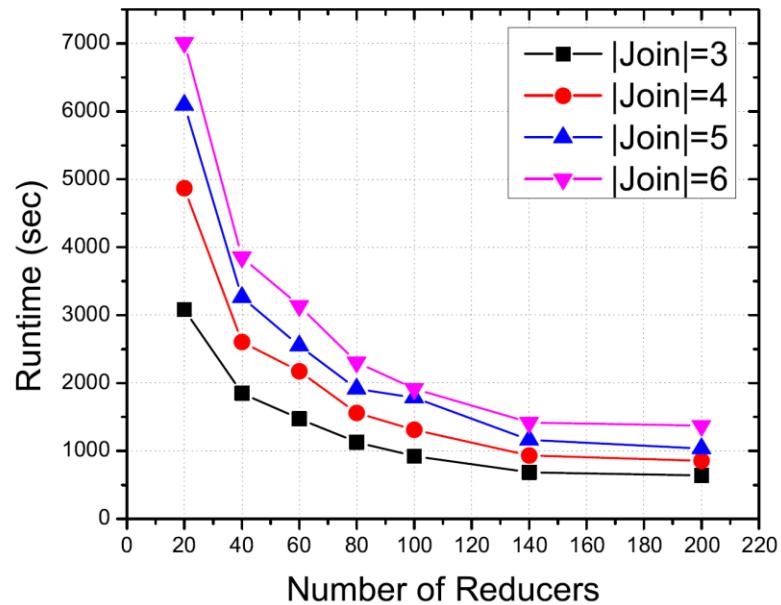


Multi-Way Spatial Join: Clique Join

R*-Tree

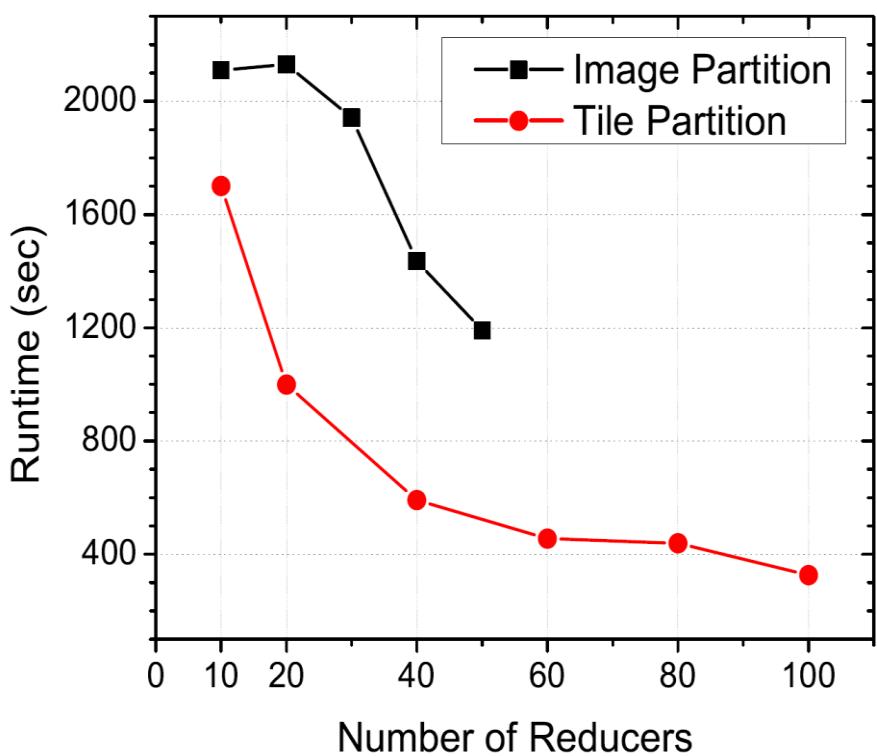


PBSM

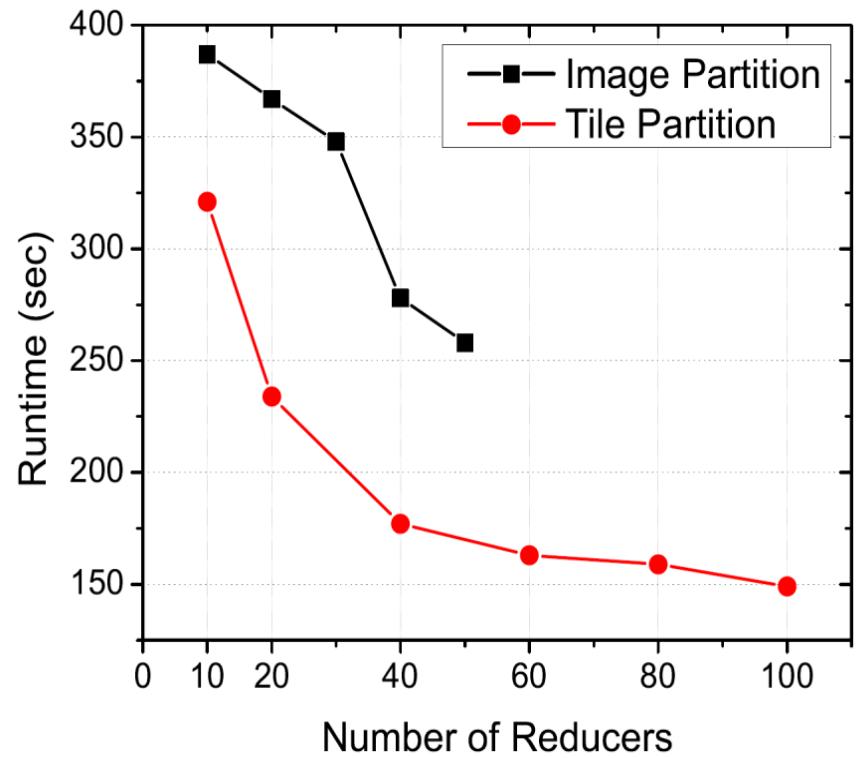


Nearest Neighbor Query

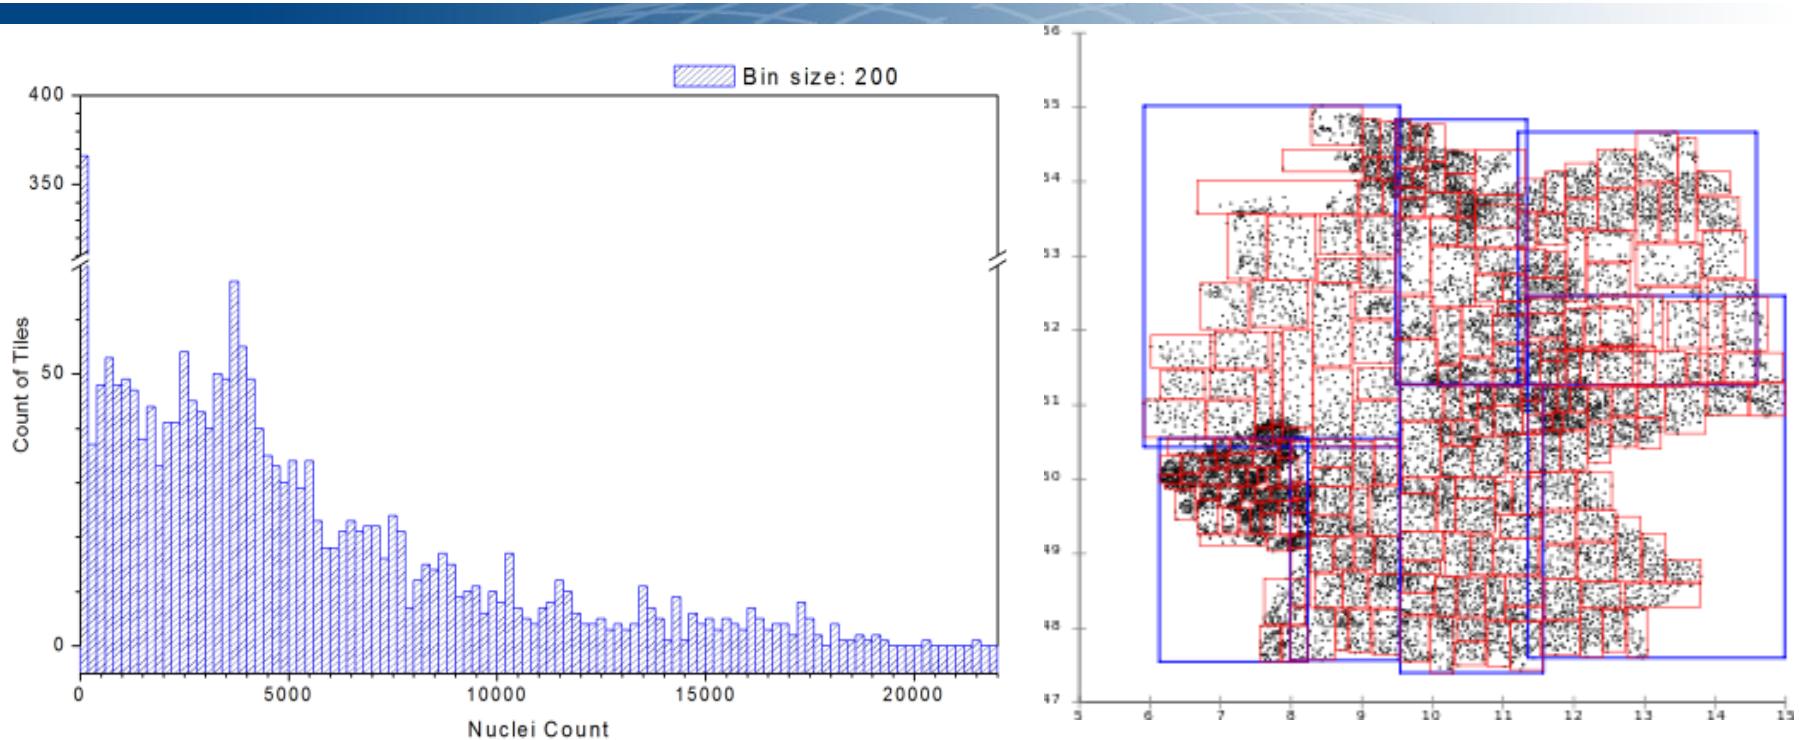
R*-Tree



Voronoi Diagram

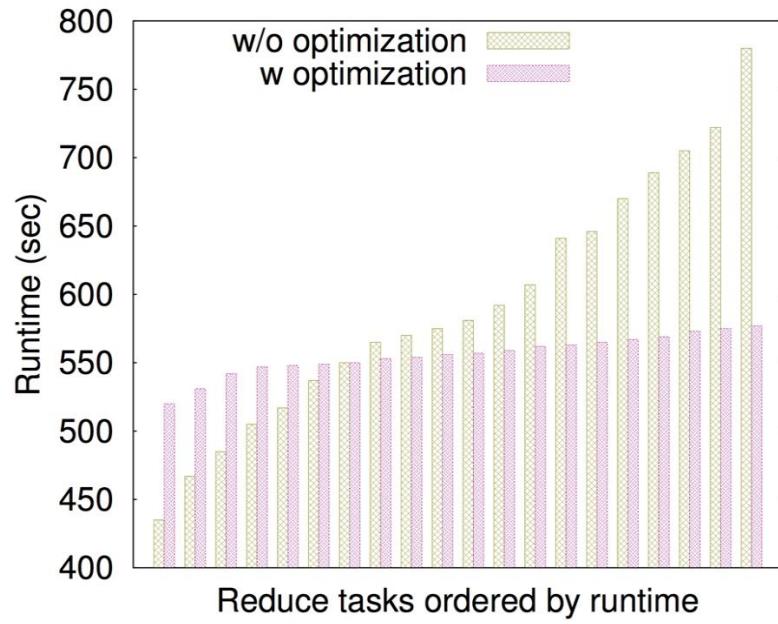
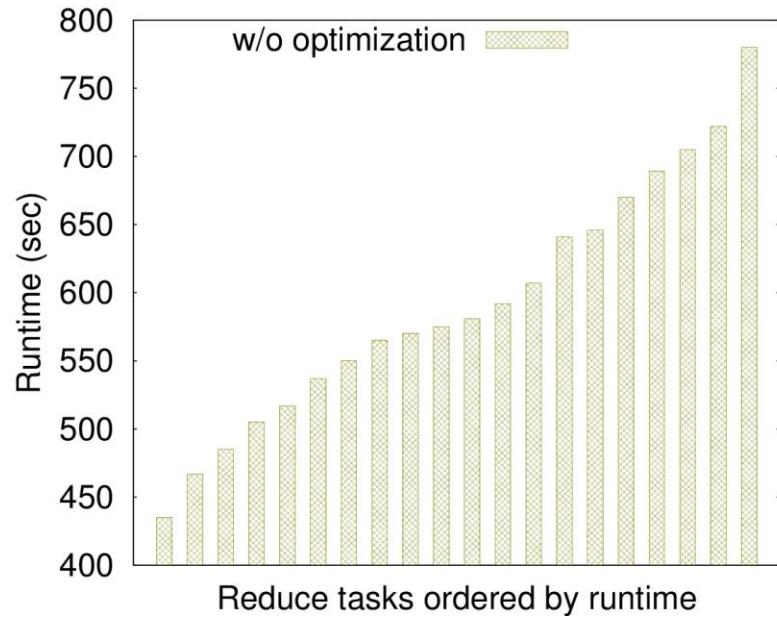


Data Skew in Spatial Query Processing



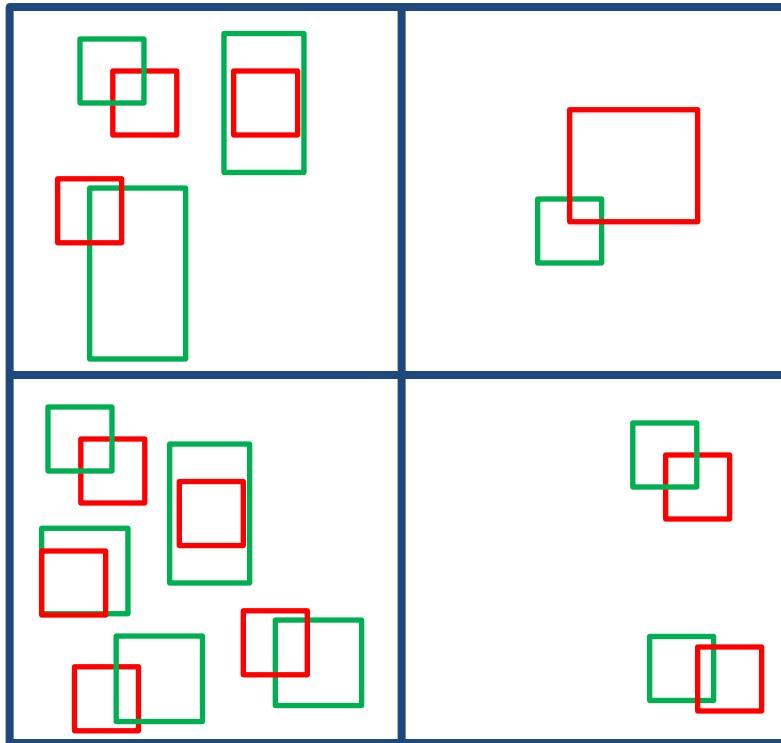
- Skew may arise for many reasons
- Longest running task becomes bottleneck

Stragglers

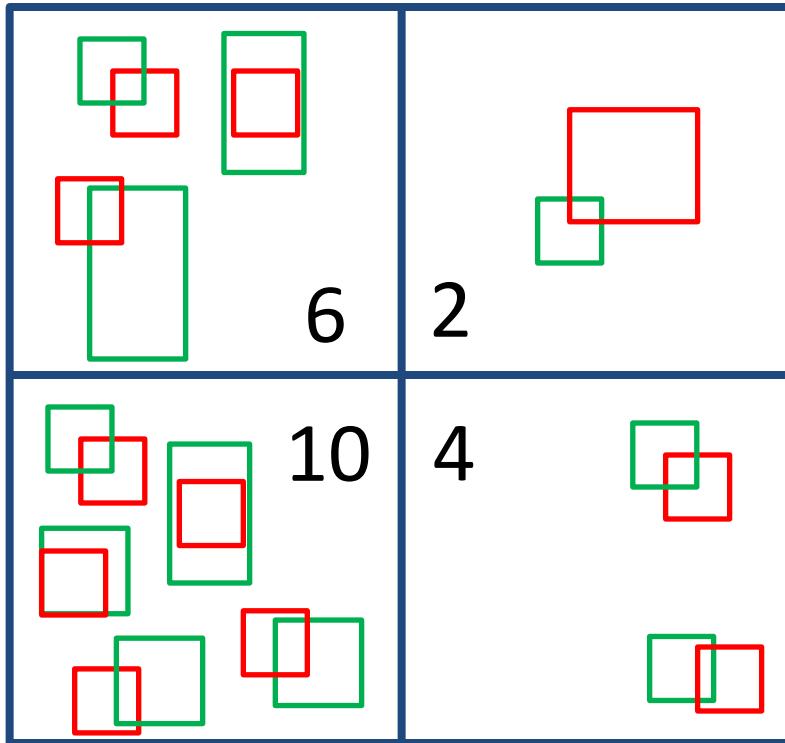


Skew Mitigation

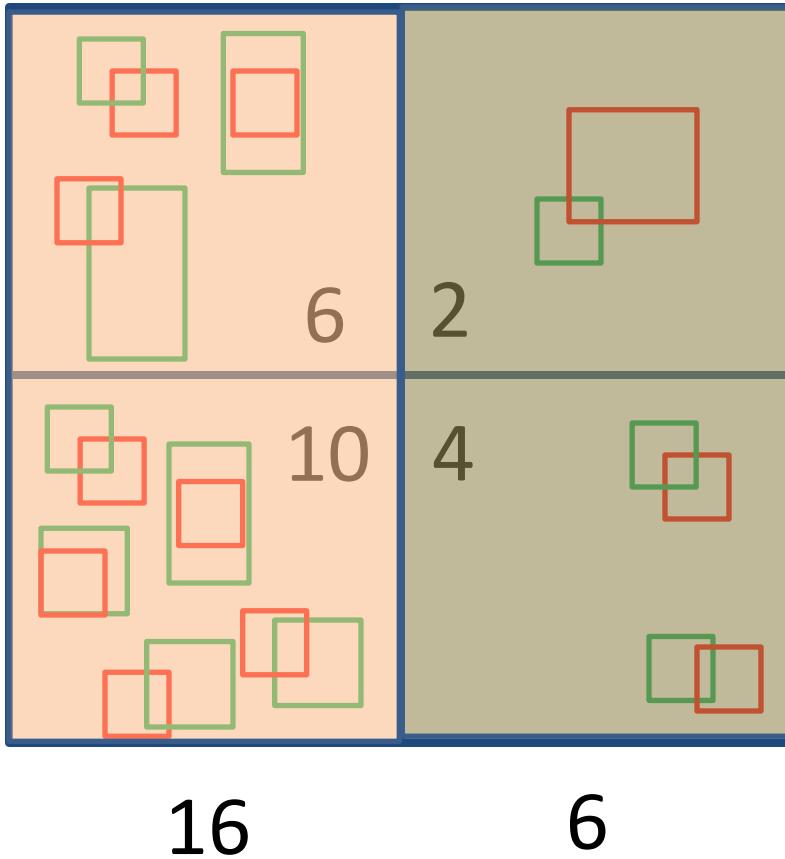
- Hadoop generates tasks using hash partitioning (default)
- Hash partitioning is not skew resistant
- We overwrite the default hashing partition function



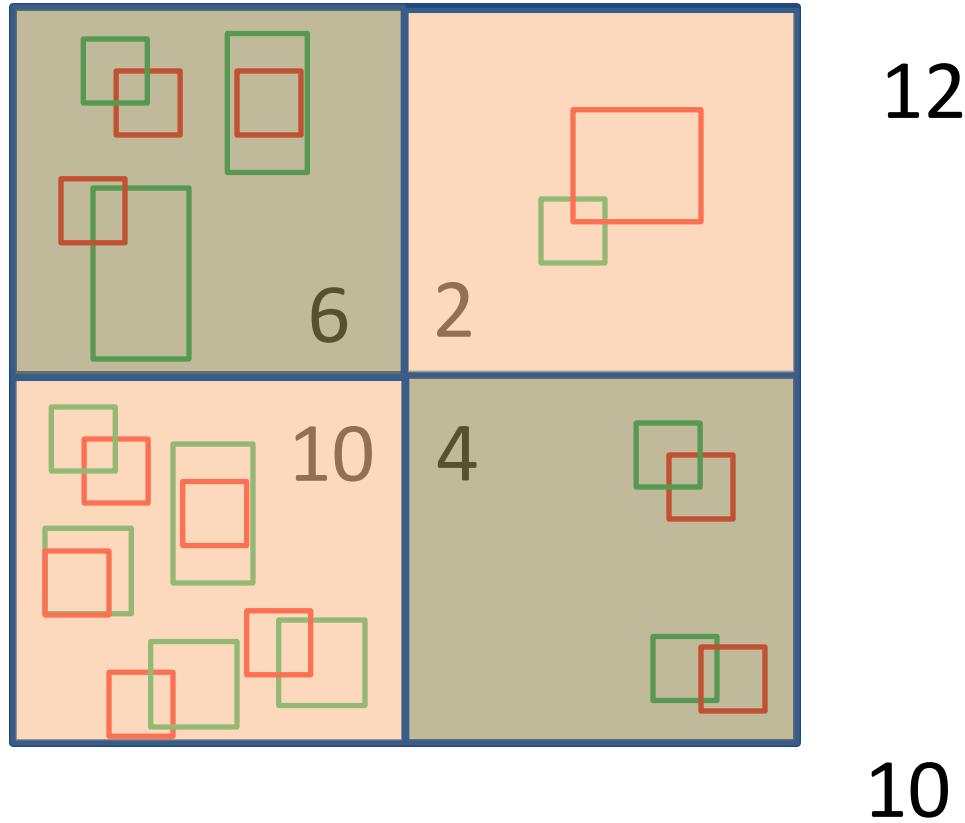
Cost Estimation



Repartition (hash)



Repartition (optimal)



Query Optimization

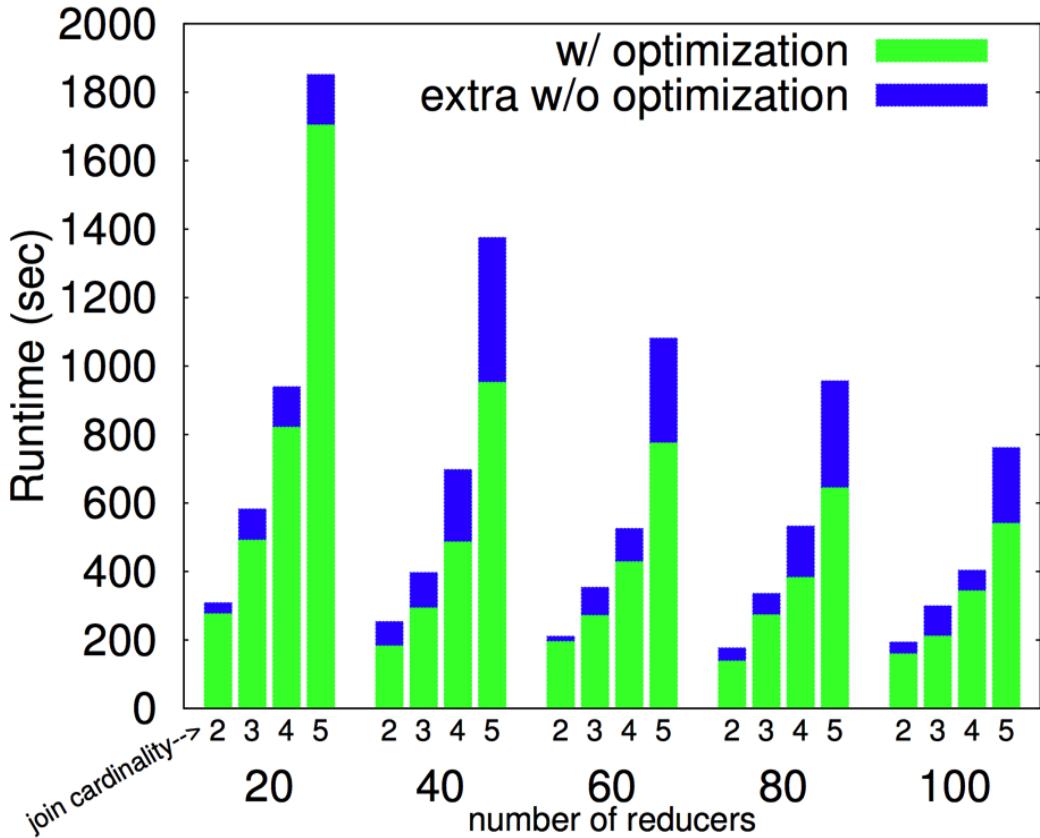
- Partition problem is NP-Complete
- Greedy Approach $\approx 4/3$ OPT (not bad)

$$Q = R \underset{i \in \text{intersects}}{\bowtie} S = \bigcup_{i=1}^N R_i \underset{i \in \text{intersects}}{\bowtie} S_i$$

$$W_j = \sum_{i \in P_j} Cost(R_i \bowtie S_i)$$

$$Cost(R_i \bowtie S_i) = \alpha|R_i| + \beta|S_i| + \gamma$$

Query Optimization



Summary

- Effective management of **big spatial data** is one of the pressing challenges for next generation integrative biomedical research.
- We propose and provide a high performance **MapReduce** based **querying system** for large scale spatial data.
- We **empirically** test the concept with analytical medical imaging as an example application.
- The system is **efficient, cost effective, scalable**, and provides **expressive query language**.

Next Steps (I)

Spatial Partitioning and Optimization Methods for MapReduce

- Density aware partition methods
 - partition according to the local spatial density
 - dynamically adjust partition granularity
- Index and pipeline selection
 - one spatial index does not fit all (R*-Tree v.s Voronoi etc..)
 - select index based on query type data set properties
 - pipeline optimization

Next Steps (II)

Spatial Query Optimization

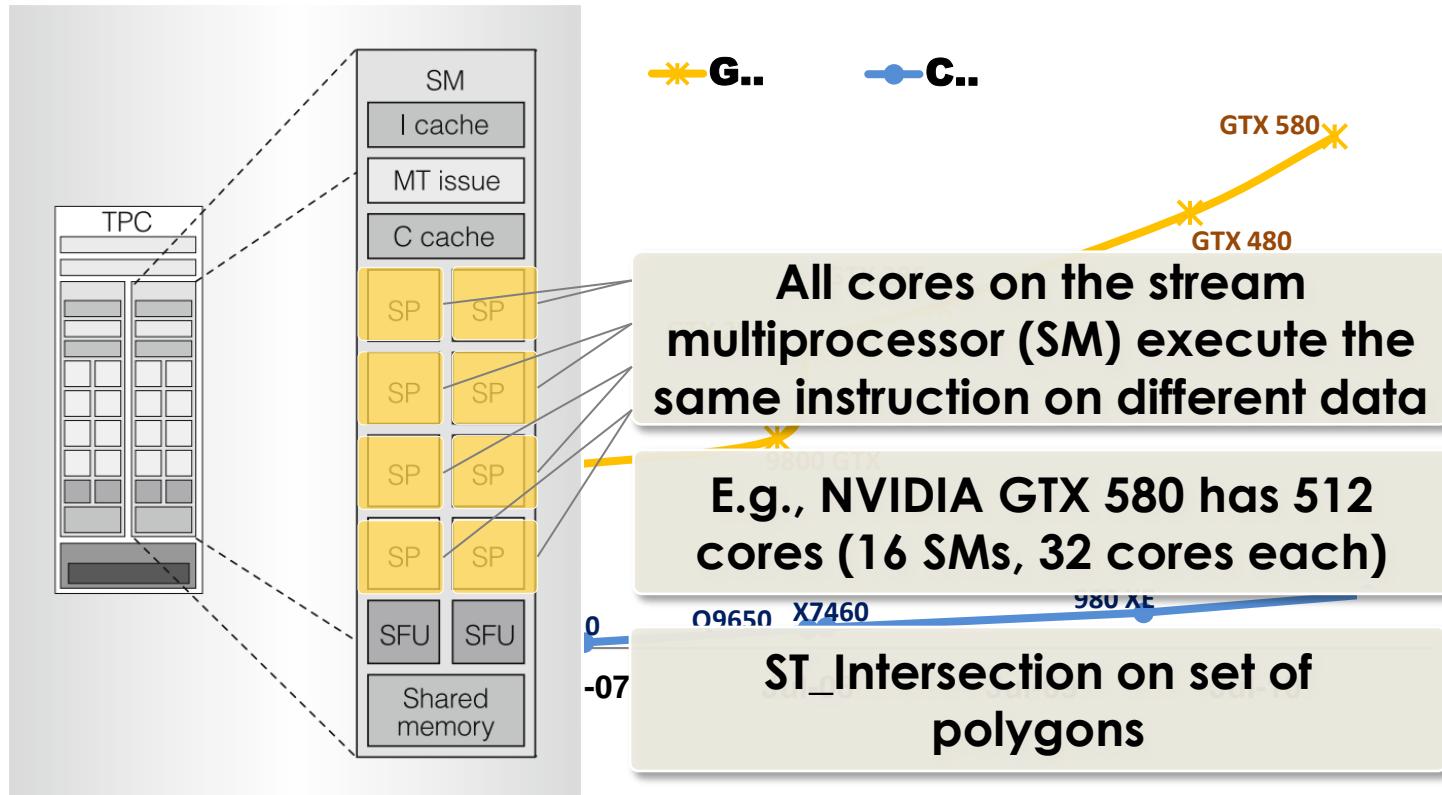
```
SELECT AVG(ratio) FROM (
    SELECT
        ST_Area(ST_Intersection(p.geom, q.geom)) /
        ST_Area(ST_Union(p.geom, q.geom)) AS ratio
    FROM    oligoastroiii_1_1 AS p, oligoastroiii_1_2 AS q
    WHERE   ST_Intersects(p.geom, q.geom) );

```

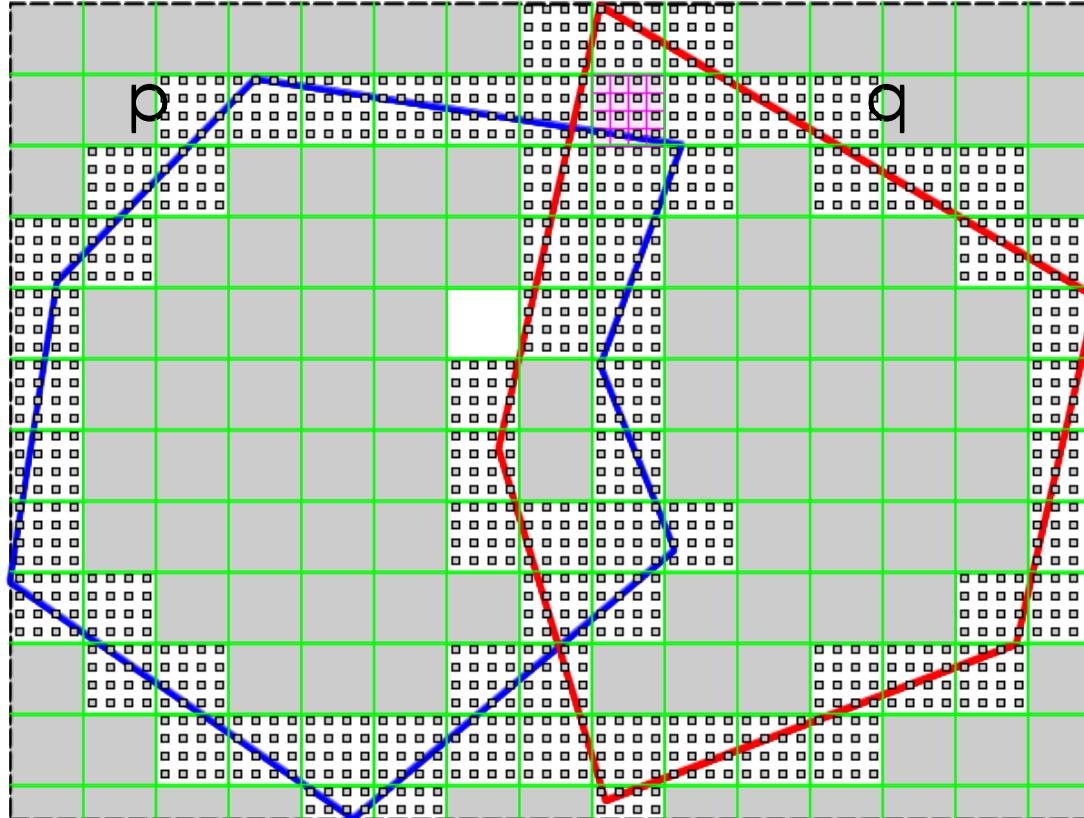
```
SELECT AVG(ratio) FROM (
    SELECT area_intersect/ (area_p + area_q -area_intersect) AS ratio
    FROM (
        SELECT ST_Area(ST_Intersection(p.geom, q.geom)) AS area_intersect,
               ST_Area(p.geom) AS area_p, ST_Area(q.geom) AS area_q
        FROM    oligoastroiii_1_1 AS p, oligoastroiii_1_2 AS q
        WHERE   p.geom && q.geom) AS temp1 )
    WHERE area_intersect >0 ) AS temp2 ;
```

Next Steps (III)

Turbo Charge Hadoop-GIS with GPU



Reduce Unnecessary Computation



K. Wang et al., Accelerating Pathology Image Data Cross-Comparison on CPU-GPU Hybrid Systems, In PVLDB 2012.

Next Steps (IV)

MapReduce + DB \approx Hybrid System

- Database
 - A mature storage engine
 - Good single node performance (40 years of research)
 - But not scalable
- MapReduce
 - Highly scalable and fault tolerant
 - But low single node throughput
- Best of the two worlds ?
 - current efforts: HadoopDB, Seongo
 - how to merge the two for big spatial data ?

Questions?

Hadoop-GIS:

<https://web.cci.emory.edu/confluence/display/HadoopGIS>